

JointAdaSpec — Joint Adaptive Speculative Decoding

Qwen 14B → 0.5B, GSM8K, McNemar paired n = 1500

+4.07% **2.22×** **3.7×**

EM vs target_only

p = 0.020 ✓

throughput vs vanilla speculative

(10.84 vs 4.88 tok/s)

faster than fuzzy_sd_T (Theorem E)

+4-8% EM advantage

Локальный заякоренный запуск 2026-05-14, RTX 5090